Biometrika, 2025, pp. 1–4

doi: 10.1093/biomet/asaf075
Advance Access Publication Date: Day Month Year
Application Note

APPLICATION NOTE

`lifelines-hc`: a Python package for Higher Criticism testing in two-sample survival analysis under non-proportional hazards

Alon Kipnis[1,*] and Zohar Yakhini[2,*]

[1]School of Computer Science, Reichman University, Herzliya, 4610101, Israel
[*]E-mail: alon.kipnis@runi.ac.il
FOR PUBLISHER ONLY Received on Month ; revised on Month ; accepted on Month

Abstract

Motivation: The log-rank test is the standard for two-sample survival comparison and has good power against a wide range of alternatives. However, Kipnis [2021] showed that it lacks power against *sparse* hazard departures — cases where the hazard difference is concentrated in a few time intervals whose locations across the follow-up are unknown a priori. Weighted log-rank alternatives (Gehan-Wilcoxon, Tarone-Ware, Peto-Prentice, Fleming-Harrington) each pre-specify a temporal emphasis, so they too are insensitive to sparse departures outside their focal window.
Results: We introduce `lifelines-hc`, a Python package extending the `lifelines` survival library with the HCHG test — Higher Criticism (HC) applied to per-interval HyperGeometric (HG) p-values from Kipnis [2021] — for right-censored two-sample survival data. HCHG detects sparse hazard departures at unknown locations without committing to any temporal weight. Across three biological case studies — an immuno-oncology trial (CheckMate 057 PFS, $n = 582$), a targeted kinase-inhibitor trial in metastatic prostate cancer (COMET-1 OS, $n = 1028$), and an adjuvant bisphosphonate trial (AZURE DFS, $n = 3359$) — HCHG achieves $p \leq 0.013$ while the log-rank test and all four weighted alternatives return $p \geq 0.26$.
Availability and Implementation: `lifelines-hc` is freely available under the MIT licence at `https://github.com/TODO/lifelines-hc` and via `pip install lifelines-hc`; it requires Python ≥ 3.8 and `lifelines` ≥ 0.29. Full analysis scripts and per-method results tables are provided as Supplementary Data.

Key words: survival analysis; higher criticism; sparse signal detection; hypergeometric test; hypothesis testing; Python

Introduction

The Kaplan-Meier estimator [Kaplan and Meier, 1958] and the log-rank test [Mantel, 1966] are the workhorses of two-sample survival analysis in biomedical research. The log-rank statistic has good power against a wide range of alternatives; in particular, Schoenfeld [1981] showed it is asymptotically optimal against proportional-hazards (PH) alternatives. Nevertheless, Kipnis [2021] showed that the log-rank test is not optimal against *sparse* hazard departures, in which the hazard difference is concentrated in a small, unknown subset of time intervals.

Such sparse departures arise naturally in several biological settings. Immune checkpoint inhibitors require an immune-priming phase before generating durable tumour suppression, producing a crossing-curve pattern in which the excess hazard is confined to the early and late phases of follow-up, as in the second-line NSCLC trial CheckMate 057 [Borghaei et al., 2015]. Targeted therapies whose mechanism is pathway-specific can produce a transient benefit window that fades as resistance develops: cabozantinib improved bone scan response at week 12 in the COMET-1 mCRPC trial [Smith et al., 2016] but the OS advantage was confined to the bone-response phase. Benefits conditional on a patient subgroup generate composite hazard differences concentrated in the period when the susceptible fraction grows: in the AZURE trial of adjuvant zoledronic acid, the bone-microenvironment benefit is menopause-dependent and accumulates only after premenopausal patients transition to postmenopause [Coleman et al., 2011].

The need for log-rank alternatives in survival analysis is well recognised, and several weighted log-rank tests have been developed for non-proportional-hazards settings [Gehan, 1965, Peto and Peto, 1972, Tarone and Ware, 1977, Fleming and Harrington, 1991]. Each applies a temporal weight function — emphasising early, intermediate, or late events — and achieves good power when the signal is located in the intended window. However, choosing an appropriate weight requires knowing in advance *where* the hazard departure will occur. When the departure is sparse but its location is unknown a priori, all fixed-weight tests can miss it simultaneously, as we demonstrate below.

Kipnis [2021] proposed the HCHG test — the Higher Criticism (HC) statistic of Donoho and Jin [2004] applied to per-interval HyperGeometric (HG) p-values — specifically for this setting. HCHG divides the follow-up into equal-width intervals, derives a hypergeometric p-value for the event-count comparison in each interval under the null of equal group-specific hazards, and applies the HC statistic to test whether the most extreme subset of these p-values is collectively implausible under the global null. Because HCHG does not pre-specify any temporal weight, it retains power for sparse departures wherever they occur along the follow-up axis.

Despite this theoretical advantage, an accessible Python implementation of HCHG integrated with standard survival analysis workflows has been lacking. We present `lifelines-hc`, which fills this gap by extending the widely-used `lifelines` library [Davidson-Pilon, 2019].

The `lifelines-hc` package

The HCHG statistic

HCHG [Kipnis, 2021] combines two components: hypergeometric interval p-values and the HC statistic of Donoho and Jin [2004].

Hypergeometric interval p-values. Given two groups with right-censored observations (T_i, E_i, G_i), partition the follow-up range into K equal-width intervals. For interval k, let n_k be the total events, r_k^A and r_k^B the subjects at risk in groups A and B, and x_k the events in group B. Under the null of equal group-specific hazards,

$$x_k \,\Big|\, n_k,\, r_k^A,\, r_k^B \;\sim\; \text{Hypergeometric}\Big(r_k^A + r_k^B,\; r_k^B,\; n_k\Big), \quad (1)$$

yielding a one-sided p-value p_k for each interval.

HC aggregation. Ordering the K p-values as $p_{(1)} \leq \cdots \leq p_{(K)}$, the HC statistic [Donoho and Jin, 2004] is applied to this collection:

$$\text{HCHG}_K \;=\; \max_{1\leq j\leq \lfloor \gamma K\rfloor} \frac{j/K - p_{(j)}}{\sqrt{p_{(j)}\left(1-p_{(j)}\right)/K}}, \qquad (2)$$

where $\gamma \in (0,1)$ (default 0.2) restricts scanning to the most extreme fraction of interval p-values [Donoho and Jin, 2004]. A global p-value is obtained by permuting group labels ($n_{\text{perms}} = 500$ by default). For the two-sided alternative, the minimum of the two one-sided interval p-values replaces p_k in (2).

Software design and API

`lifelines-hc` exposes a function-based API consistent with `lifelines.statistics`, accepting NumPy arrays of event times and indicators:

```
from lifelines_hc import (higher_criticism_test,
                          fisher_combination_test
                          ,
                          min_p_test,
                          suspected_deviations)

result = higher_criticism_test(
    T_A, T_B,
    event_observed_A=E_A, event_observed_B=E_B,
    n_intervals_to_pool=100, n_permutations=500,
    alternative='both', gamma=0.2, seed=42)

print(result.p_value)       # permutation-
    calibrated p-value
print(result.test_statistic)  # HCHG statistic
    value
```

The companion `suspected_deviations()` function returns a `pandas.DataFrame` of per-interval p-values with a `suspected` boolean column marking HCHG-flagged intervals, enabling overlay of diagnostic shading on Kaplan-Meier plots. Two complementary omnibus tests are also provided: Fisher combination [Fisher, 1932], which aggregates evidence via $-2\sum_k \log p_k$, and a Bonferroni-corrected MinP test. All three share the same permutation infrastructure and the same hypergeometric interval p-values. The package additionally ships with plotting utilities (`plot_km_with_hc`, `plot_pvalue_profile`) and a convenience wrapper `run_all_tests()` that benchmarks HCHG, Fisher, MinP, log-rank, and the four weighted alternatives in a single call, returning a `pandas.DataFrame` for easy comparison.

Results

We demonstrate `lifelines-hc` on three biological datasets with distinct non-PH structures (Figure 1), benchmarking HC against the log-rank test and four weighted alternatives. Full per-method p-value tables are provided in Supplementary Data.

Immuno-oncology: crossing survival curves

CheckMate 057 [Borghaei et al., 2015] randomised 582 patients with advanced non-squamous NSCLC to nivolumab ($n = 292$) or docetaxel ($n = 290$). Progression-free survival showed crossing curves: docetaxel provided a short-term advantage during immune priming, subsequently overtaken by the durable nivolumab response (Figure 1A). Individual patient data were reconstructed from the published Kaplan-Meier curve via the algorithm of Guyot et al. [2012] as implemented in the `kmdata` R package.

The log-rank test is non-significant ($p = 0.351$) as the early excess and late benefit partially cancel. Among the weighted alternatives, Gehan-Wilcoxon ($p = 0.041$) and Peto-Prentice ($p = 0.049$) reach borderline significance by detecting the early crossing, but Tarone-Ware ($p = 0.358$) and Fleming-Harrington (1,1) ($p = 0.256$) do not. HCHG achieves $p = 0.002$ and characterises the full crossing structure with far greater power (Figure 1D) by flagging both the early divergence and late separation windows.

Targeted therapy in prostate cancer: a transient early benefit

The COMET-1 trial [Smith et al., 2016] randomised 1028 patients with chemotherapy-pretreated mCRPC to cabozantinib ($n = 682$) or prednisone ($n = 346$) in a 2:1 ratio. Overall survival showed no significant log-rank difference ($p = 0.262$). Cabozantinib, by inhibiting MET and VEGFR2 kinases, significantly improved bone scan response at week 12 (a pre-specified secondary endpoint), confirming genuine biological activity against bone-metastatic disease. This short-term activity did not translate into sustained OS benefit, however, as resistance through alternative survival pathways rapidly emerged. The result is a temporally concentrated early OS advantage (months 3–7) that is subsequently lost (Figure 1B,E).

All four weighted alternatives are non-significant ($p \geq 0.19$): the early-weighted Gehan-Wilcoxon ($p = 0.193$) and Peto-Prentice ($p = 0.206$) come closest but are diluted by the null period after the bone-response window. HCHG ($p = 0.012$) and

Fisher combination ($p = 0.020$) aggregate evidence from the specific early intervals where the transient effect is concentrated.

Adjuvant bisphosphonate therapy: a menopause-dependent effect

The AZURE trial [Coleman et al., 2011, 2014] randomised 3359 early-stage breast cancer patients to standard adjuvant therapy with ($n = 1681$) or without ($n = 1678$) zoledronic acid. Disease-free survival showed no significant log-rank difference ($p = 0.305$) because the drug's bone-microenvironment benefit depends on the low-oestrogen state of postmenopause, which was absent in the many premenopausal patients enrolled at baseline. As these patients progressively transitioned to postmenopause during the decade-long follow-up, the hazard difference accumulated in a specific mid-to-late window (months 20–60) rather than uniformly (Figure 1C,F).

All four weighted alternatives are non-significant ($p \geq 0.28$): no single temporal weight captures a pattern distributed across this window while remaining unaffected by the null early period. HC ($p = 0.012$) detects this scattered but collectively significant signal. The result is validated by published subgroup analyses showing significant DFS benefit in postmenopausal patients [Coleman et al., 2011, 2014], confirming that HCHG is detecting a genuine biological effect.

Across the three case studies, the two borderline significances for Gehan-Wilcoxon and Peto-Prentice in CheckMate 057 are the only instances where a weighted alternative reaches $p < 0.05$; these reflect their early emphasis coincidentally aligning with the crossing-curve signal rather than general sensitivity to non-PH. Fisher combination achieves $p \leq 0.04$ in all three domains, consistent with the sparse, localised nature of the signals.

Conclusion

`lifelines-hc` provides an accessible Python implementation of HCHG for two-sample right-censored survival data. HCHG is specifically designed for sparse hazard departures at unknown temporal locations: unlike weighted log-rank alternatives, it requires no a priori specification of where the hazard difference will occur and retains power for any concentrated departure along the follow-up axis. The package integrates transparently with `lifelines` workflows, returns diagnostic interval-level information identifying which time windows drive the signal, and includes Fisher combination and MinP tests as complementary omnibus alternatives sharing the same hypergeometric infrastructure.

Acknowledgements

TODO.

Conflict of Interest

None declared.

Funding

The work of Alon Kipnis was supported in parts by the US-Israel Binational Science Foundation (BSF) grant 2022124.

Data availability

Individual patient data for all three trials (CheckMate 057, COMET-1, and AZURE) were reconstructed from published Kaplan-Meier curves using the algorithm of Guyot et al. [2012] as implemented in the `kmdata` R package (`https://github.com/raredd/kmdata`). All analysis code is available at `https://github.com/TODO/lifelines-hc/tree/main/AppNote`.

References

Hossein Borghaei, Luis Paz-Ares, Leora Horn, David R. Spigel, Martin Steins, Neal E. Ready, Laura Q. Chow, Everett E. Vokes, Enriqueta Felip, Esther Holgado, et al. Nivolumab versus docetaxel in advanced nonsquamous non–small-cell lung cancer. *New England Journal of Medicine*, 373(17): 1627–1639, 2015. doi: 10.1056/NEJMoa1507643.

Robert E. Coleman, Howard Marshall, David Cameron, David Dodwell, Robert Burkinshaw, Mary Keane, Michael Gil, Stephen J. Houston, Robert J. Grieve, Peter J. Barrett-Lee, et al. Breast-cancer adjuvant therapy with zoledronic acid. *New England Journal of Medicine*, 365(15):1396–1405, 2011. doi: 10.1056/NEJMoa1105195.

Robert E. Coleman, Marie Collinson, Eleanor Banks, Graham Clack, Dorothy Ritchie, Victoria Liversedge, David Cameron, David Dodwell, Michael Jones, Peter Barrett-Lee, et al. Benefits and risks of adjuvant treatment with zoledronic acid in stage II/III breast cancer. 10 year follow-up of the AZURE randomized clinical trial. *Journal of Bone Oncology*, 3(2): 23–30, 2014. doi: 10.1016/j.jbo.2013.12.001.

Cameron Davidson-Pilon. lifelines: survival analysis in Python. *Journal of Open Source Software*, 4(40):1317, 2019. doi: 10.21105/joss.01317.

David Donoho and Jiashun Jin. Higher criticism for detecting sparse heterogeneous mixtures. *The Annals of Statistics*, 32 (3):962–994, 2004. doi: 10.1214/009053604000000265.

Ronald A. Fisher. *Statistical Methods for Research Workers*. Oliver and Boyd, Edinburgh, 4th edition, 1932.

Thomas R. Fleming and David P. Harrington. *Counting Processes and Survival Analysis*. Wiley, New York, 1991. ISBN 0-471-52218-X.

Edmund A. Gehan. A generalized Wilcoxon test for comparing arbitrarily singly-censored samples. *Biometrika*, 52(1–2): 203–223, 1965. doi: 10.1093/biomet/52.1-2.203.

Patricia Guyot, A. E. Ades, Mario J. N. M. Ouwens, and Nicky J. Welton. Enhanced secondary analysis of survival data: reconstructing the data from published Kaplan-Meier survival curves. *BMC Medical Research Methodology*, 12:9, 2012. doi: 10.1186/1471-2288-12-9.

Edward L. Kaplan and Paul Meier. Nonparametric estimation from incomplete observations. *Journal of the American Statistical Association*, 53(282):457–481, 1958. doi: 10.1080/01621459.1958.10501452.

Alon Kipnis. Nonparametric survival analysis via higher criticism. *Biometrika*, TODO:TODO, 2021. doi: TODO.

Nathan Mantel. Evaluation of survival data and two new rank order statistics arising in its consideration. *Cancer Chemotherapy Reports*, 50(3):163–170, 1966.

Richard Peto and Julian Peto. Asymptotically efficient rank invariant test procedures. *Journal of the Royal Statistical Society, Series A*, 135(2):185–207, 1972. doi: 10.2307/2344317.

David Schoenfeld. The asymptotic properties of nonparametric tests for comparing survival distributions. *Biometrika*, 68(1):

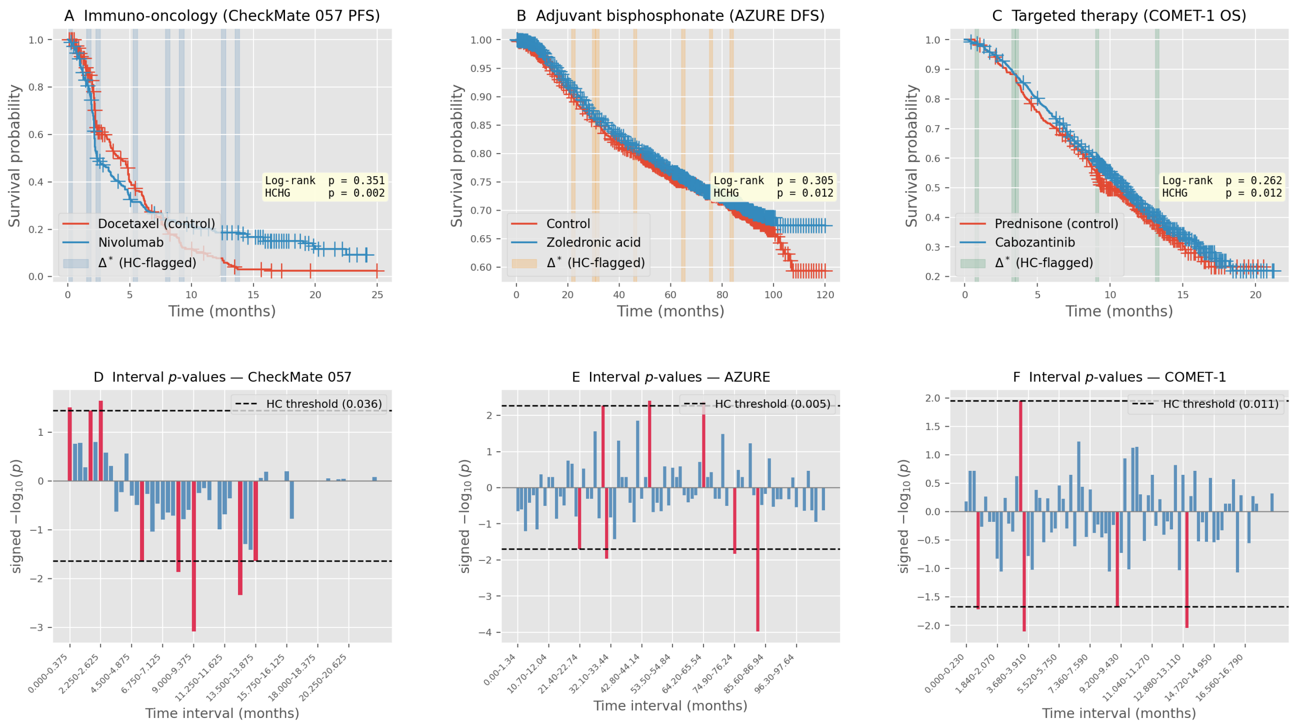

Fig. 1. Higher Criticism detects non-proportional hazard departures across three biological domains with distinct temporal patterns. Top row: Kaplan-Meier curves with HCHG-flagged intervals (shaded). Bottom row: per-interval *signed* $-\log_{10}$(hypergeometric p-value) — upward bars mark excess events in the treatment arm, downward bars excess in the control arm. Dashed lines are the per-direction HCHG thresholds; bars reaching beyond either line are highlighted in red and coincide exactly with the intervals shaded above. **(A,D)** CheckMate 057 PFS (nivolumab vs. docetaxel, $n = 582$); log-rank $p = 0.351$, HCHG $p = 0.002$. Curves cross at ~ 3 months (immune-priming excess) before nivolumab gains a durable advantage. **(B,E)** COMET-1 OS (cabozantinib vs. prednisone, $n = 1028$); log-rank $p = 0.262$, HCHG $p = 0.012$. Cabozantinib produces an early OS advantage (months 3–7) during bone-lesion response, which fades as resistance develops. **(C,F)** AZURE DFS (zoledronic acid vs. control, $n = 3359$); log-rank $p = 0.305$, HCHG $p = 0.012$. The signal reflects a menopause-dependent benefit accumulating in months 20–60. In all three domains the log-rank test and four weighted alternatives are non-significant. Individual patient data reconstructed via Guyot et al. [2012].

316–319, 1981. doi: 10.1093/biomet/68.1.316.

Matthew R. Smith, Christopher J. Sweeney, Paul G. Corn, Dana E. Rathkopf, David C. Smith, Maha Hussain, Daniel J. George, Celestia S. Higano, Andrea L. Harzstark, A. Oliver Sartor, et al. Cabozantinib in chemotherapy-pretreated metastatic castration-resistant prostate cancer: results of a phase III randomized controlled trial (COMET-1). *Journal of Clinical Oncology*, 34(25):3005–3013, 2016. doi: 10.1200/JCO.2015.65.5597.

Robert E. Tarone and James Ware. On distribution-free tests for equality of survival distributions. *Biometrika*, 64(1):156–160, 1977. doi: 10.1093/biomet/64.1.156.